

Deep Dive into Container Networking v1.0

Speakers:

- Nguyen Phuong An (<u>AnNP@vn.fujitsu.com</u>)
- Kim Bao Long (longkb@vn.fujitsu.com)
- Nguyen Hai Truong (<u>truongnh@vn.fujitsu.com</u>)



Agenda



- \$ whoarewe
- Container networking 101
 - Single host
 - Multiple host
 - Orchestration
 - Services discovery
 - Container network interface
- Kubernetes networking

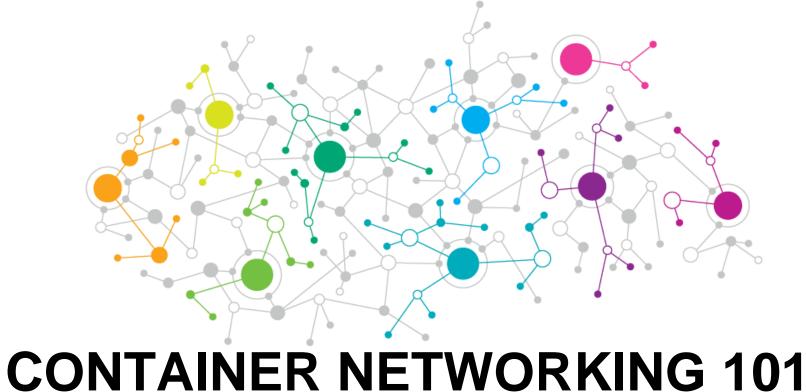
\$ whoarewe



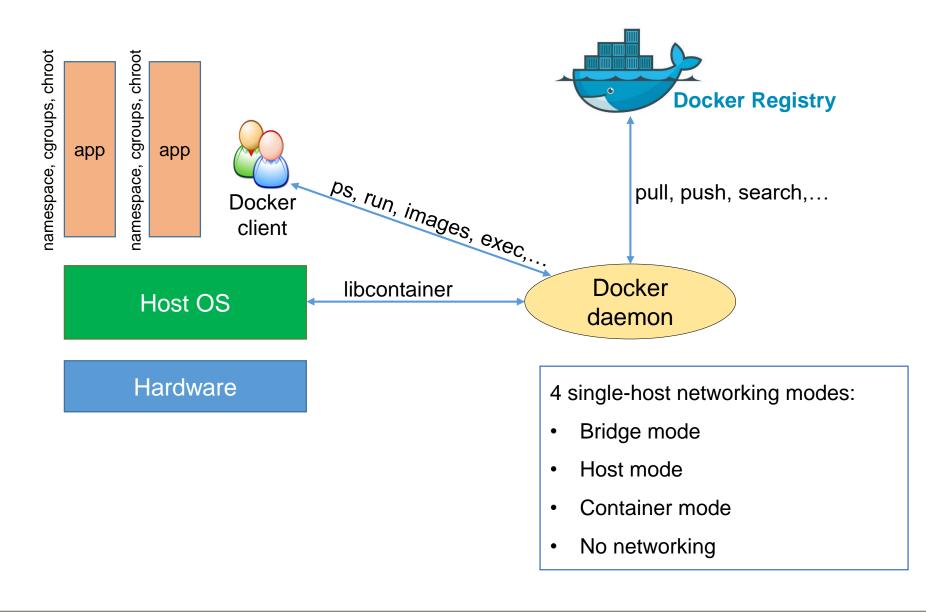
- Software engineers
- 3 years experience contribute to OpenStack Networking
 - 116 commits and 26960 LOCs
 - Neutron packet logging, firewall as a services v2.0, deploy Neutron-api under WSGI server,...
- Now, we're moving to Kubernetes
- Twitter: @annp87, @long_kb, @truongnh92





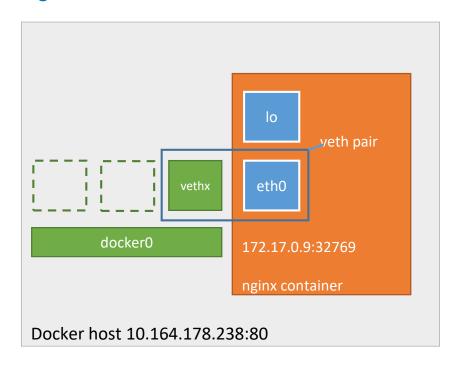








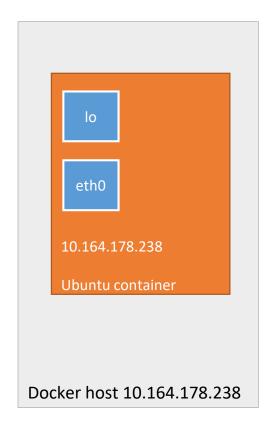
Bridge Mode Networking



```
$ docker ps
CONTAINER ID IMAGE COMMAND CREATED
STATUS PORTS NAMES
cecf17dc9117 nginx:1.9.1 "nginx -g 'daemon ..." 5 minutes ago Up 5 minutes
0.0.0.0:32769->80/tcp,
0.0.0.0:32768->443/tcp loving_engelbart
```



Host Mode Networking



- Disable network isolation of a Docker container
- The container shares the network namespace of the host
- No routing overhead → faster than bridge mode
- Expose the container directly to public network → security implications



Container Mode Networking

Docker reuse the network namespace of another container

```
$ docker run -d -P --net=bridge nginx:1.9.1
$ docker ps
CONTAINER ID
                  IMAGE
                                COMMAND
                                                                     STATUS
                                                      CREATED
                                                                                        PORTS
                                                                                                       NAMES
                                 "nginx -g 'daemon ..." 5 minutes ago
cecf17dc9117
                 nginx:1.9.1
                                                                        Up 5 minutes
0.0.0.0:32769->80/tcp, 0.0.0.0:32768->443/tcp loving engelbart
$ docker exec -it loving engelbart ip addr
36: eth0@if37: <BROADCAST,MULTICAST,UP,LOWER UP> mtu 1500 qdisc noqueue state UP group default
  link/ether 02:42:ac:11:00:09 brd ff:ff:ff:ff:ff
  inet 172.17.0.9/16 scope global eth0
   valid Ift forever preferred Ift forever
$ docker run -it --net=container:loving engelbart ubuntu:14.04 ip addr
36: eth0@if37: <BROADCAST,MULTICAST,UP,LOWER UP> mtu 1500 qdisc noqueue state UP group default
  link/ether 02:42:ac:11:00:09 brd ff:ff:ff:ff:ff
  inet 172.17.0.9/16 scope global eth0
   valid Ift forever preferred Ift forever
  inet6 fe80::42:acff:fe11:9/64 scope link
   valid Ift forever preferred Ift forever
```



No Networking

There is no networking configured

```
$ docker run -d -P --net=none nginx:1.9.1
$ docker ps
CONTAINER ID
                 IMAGE
                                COMMAND
                                                    CREATED
                                                                   STATUS
                                                                                  PORTS
                                                                                                               NAMES
                                "nginx -g 'daemon ..." About a minute ago Up About a minute
df7df16a84de
                 nginx:1.9.1
                                                                                                             friendly_northcutt
$ docker inspect df7df16a84de | grep IPAddress
"SecondaryIPAddresses": null,
      "IPAddress": "",
          "IPAddress": "",
```



When scaling out horizontally a cluster of machines:

- How do containers talk to each other on different hosts?
- Communication between containers and with the outside world?

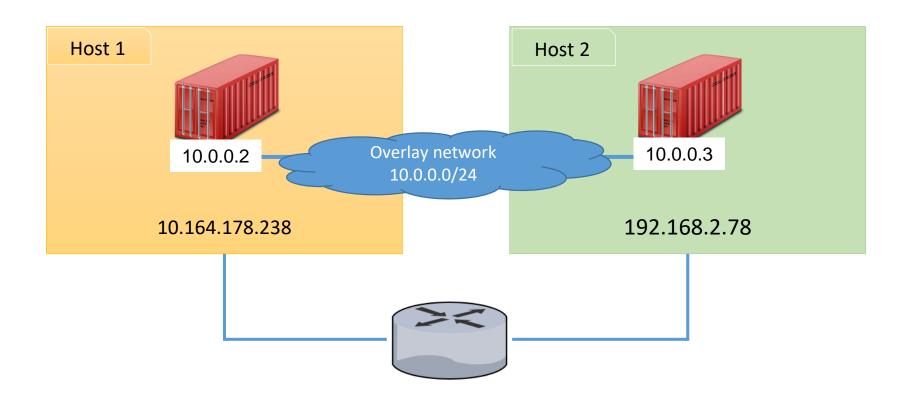
Overlay Networks

Network plugins:

- Flannel
- Calico
- Weave Net
- ...



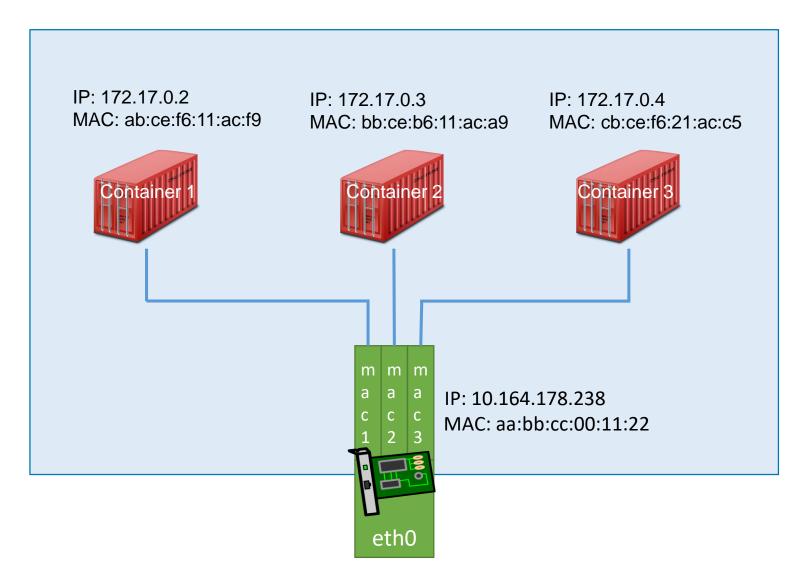
Overlay Networking



Creating a distributed network across hosts on top of the host-specific network

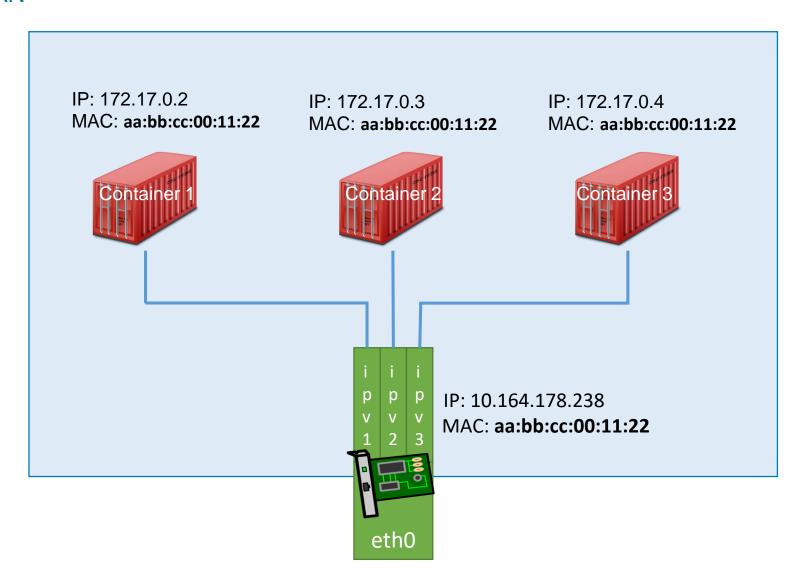


MACVLAN



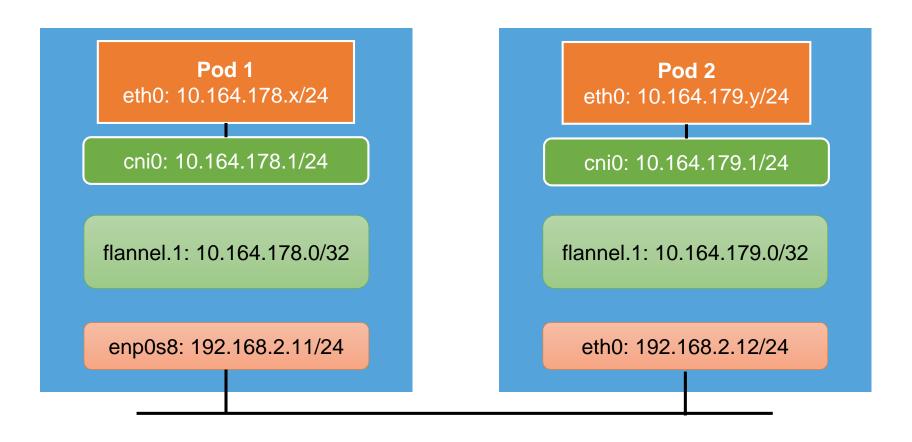


IPVLAN



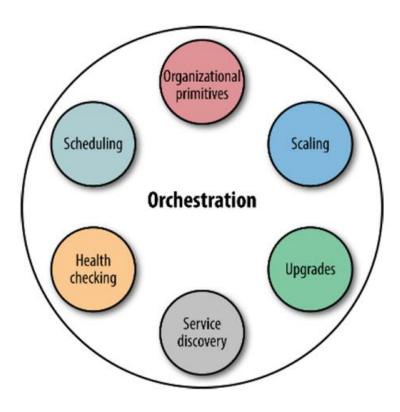


flannel



Orchestration





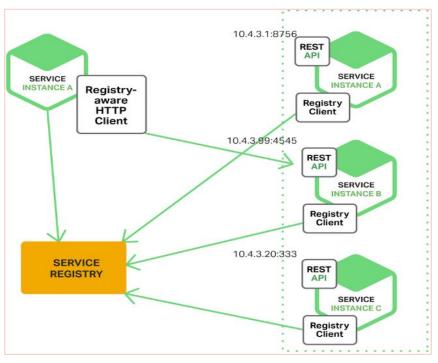
Service discovery

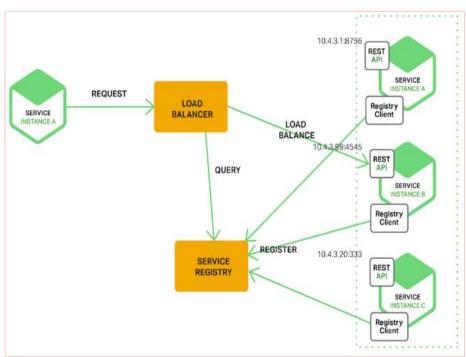


- How do you determine which host your container ended up being scheduled on so that you can connect to it?
 - maintaining a mapping between a running container and its location
 - 2 operations must be supported by a container service discovery solution
 - Registration: Establishes the container -> location mapping
 - Lookup: Enables other services or applications to look up the mapping we stored during registration

Service Discovery pattern







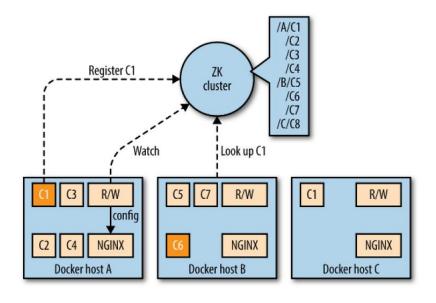
Client-side discovery pattern

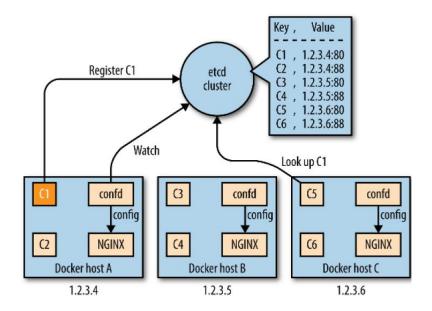
Server-side discovery pattern

Services discovery technology



- Zookeeper
- Etcd
- Consul,
- **...**





Container network interface

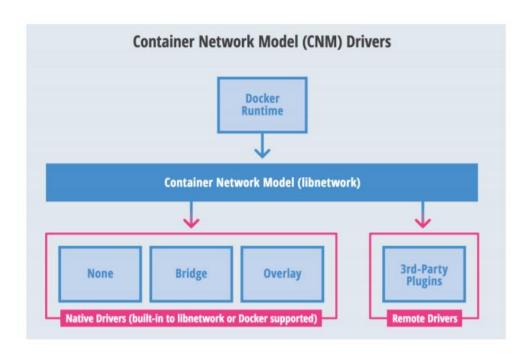


There are two proposed standards for configuring network interfaces for Linux containers: CNM & CNI

Container network model (CNM)

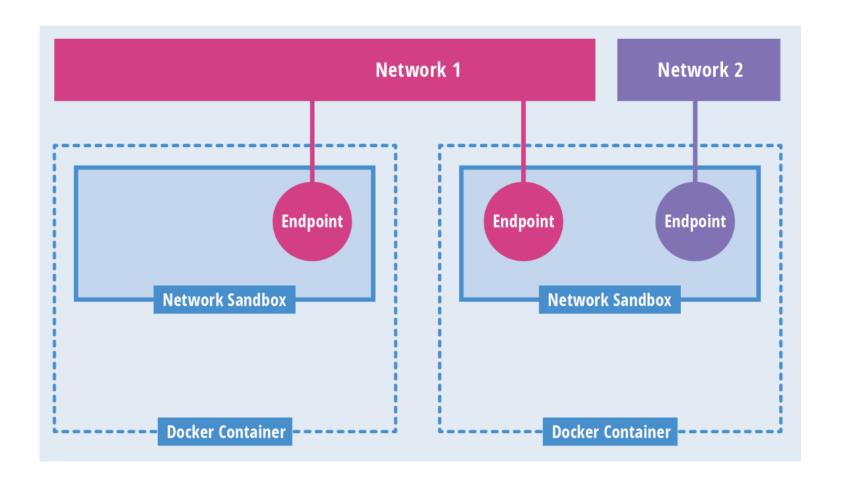


- Specification proposed by Docker and adopted by libnetwork
- Supports only Docker runtime
- Integrate with any kind of networking technology to connect and discover containers



Container network model





CNM topology

Container network model





Get/release Pool

Get/release IP

Create/delete Network

Create/delete Endpoint

Join/leave Endpoint

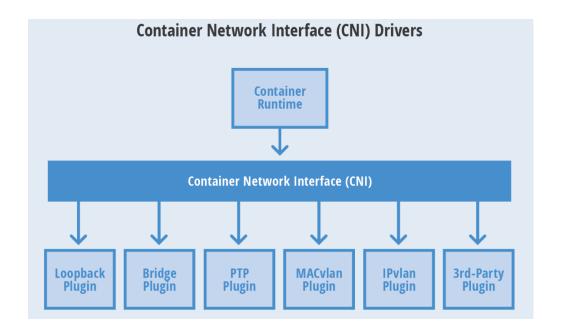
CNM IPAM

CNM Network

Container network interface



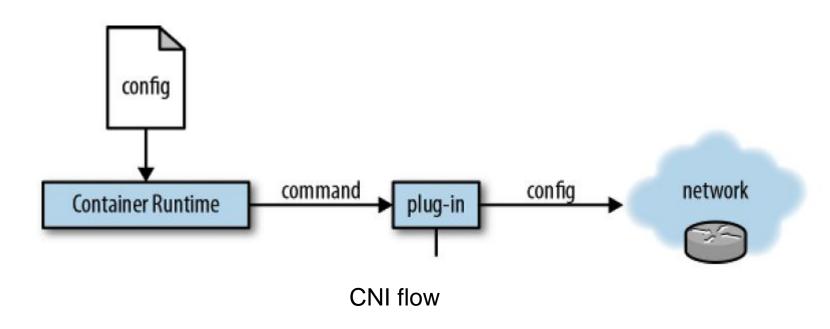
- Specification proposed by CoreOS and adopted by projects such as Kubernetes
- Supports any container runtime
- Only dealing with connectivity between containers



Container network interface



- CNI enable you to do add/remove network interfaces
- CNI defines the following operations
 - Add container to one or more networks
 - Delete container from network
 - Report CNI version



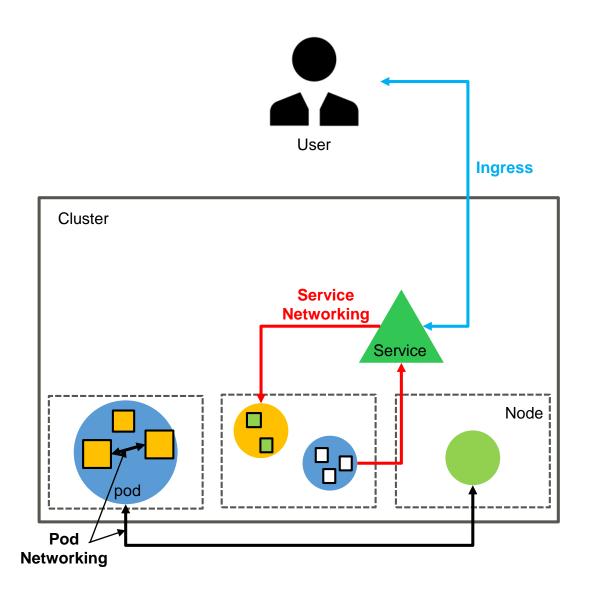


KUBERNETES NETWORKING

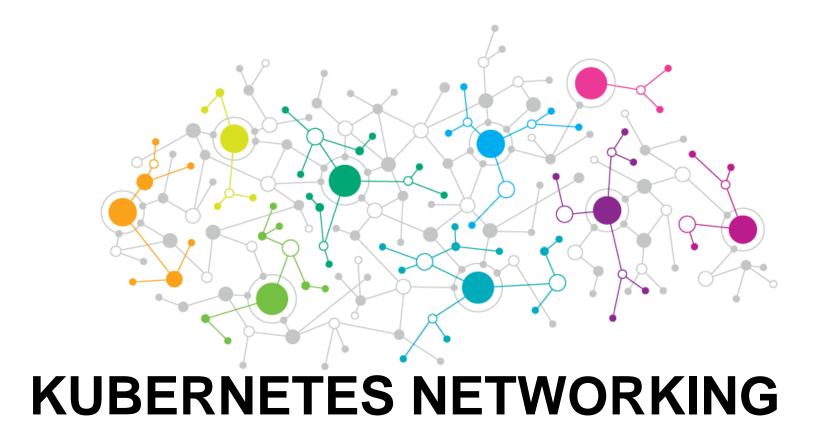
Kubernetes Networking



- Pod Networking
- Service Networking
- Ingress



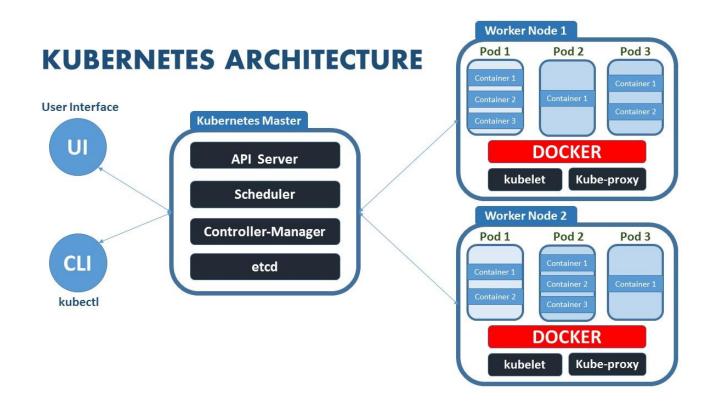




Fundamental Requirements



- Kubernetes only states three fundamental requirements
 - Containers can communicate with all others without NAT
 - Nodes can communicate with all containers without NAT
 - The IP a container sees itself is the same IP as others see it





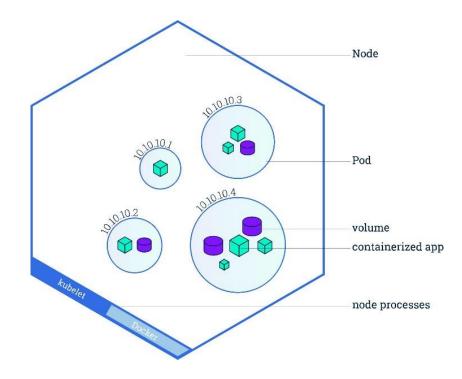
The network that enables pods to connect to each other across nodes in a kubernetes cluster

POD NETWORKING

Kubernetes Pod



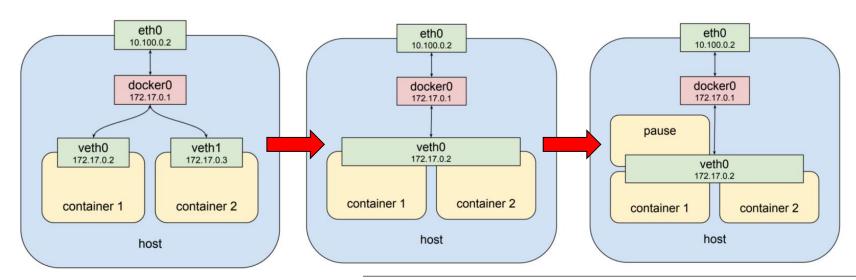
- What is Pod?
 - Basic unit in K8s
 - One or more containers that are collocated on the same host, and are configured to share a network stack and other resources (volume)



From Docker to K8s networking



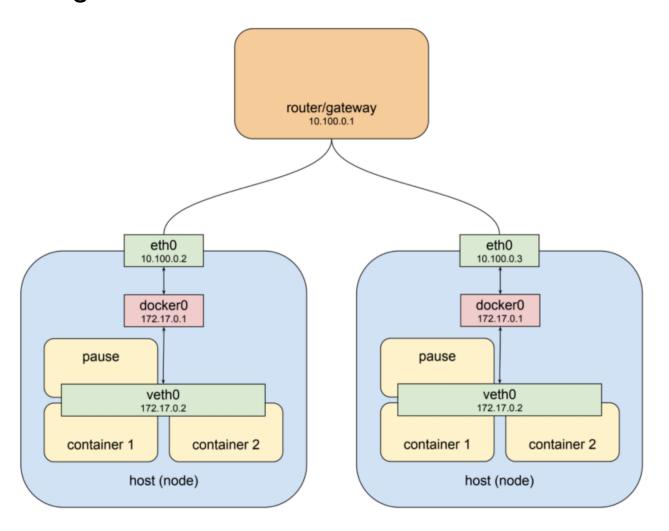
- Docker containers on local machine
 - Communicate via docker0 bridge
 - Not "share network stack"
- => Share an existing interface
 - Both containers are addressable from the outside (172.17.0.2)
 - The inside each can hit ports opened by the other on localhost
 - Restriction: cannot open the same port (same with multiple processes)
- In Kubernetes
 - Using pause container as a heart of the pod



Pod networking



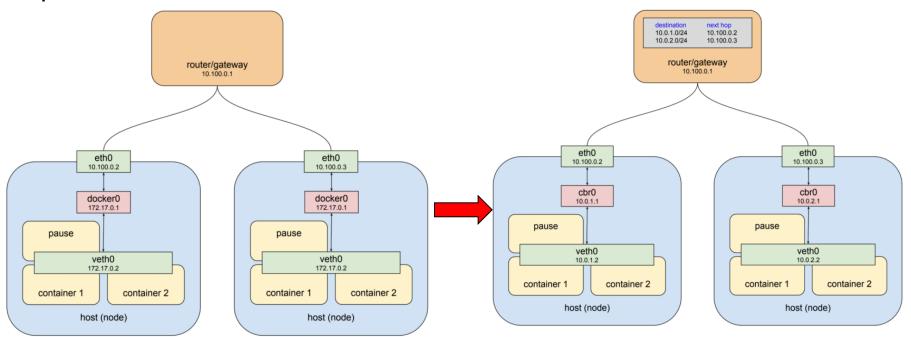
Requirement: Pods are able to communicate with the others in both single host and multi-host without NAT



Pod networking



- Each pod has a routable IP
- Kubernetes built an overlay network by:
 - Assigns an overall address space for the bridges on each node
 - Assigns the bridges addresses within that space, based on the node the bridge is built on
 - Adds routing rules to the gateway telling it how packets destined for each bridge should be routed
- pods can communicate without NAT





How the service network provides load balancing for pods so that clients inside the cluster can communicate with them reliably

SERVICE NETWORKING

Kubernetes Service

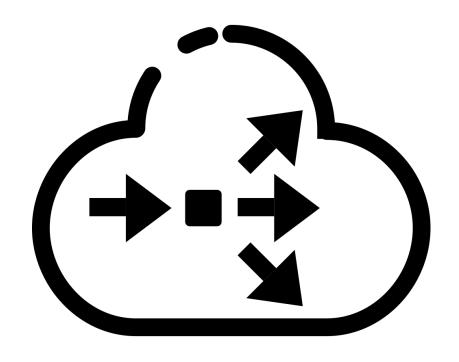


- Problem: pods in kubernetes are ephemeral
 - Auto-scaling
 - Auto-healing



Address changing

- Migration
- => Standard solution
 - Clients connect to the proxy
 - The proxy responsibilities:
 - Durable and resistant to failure
 - Maintain a list of healthy servers
 - Having a keep alive mechanism
- => Service



Service Networking

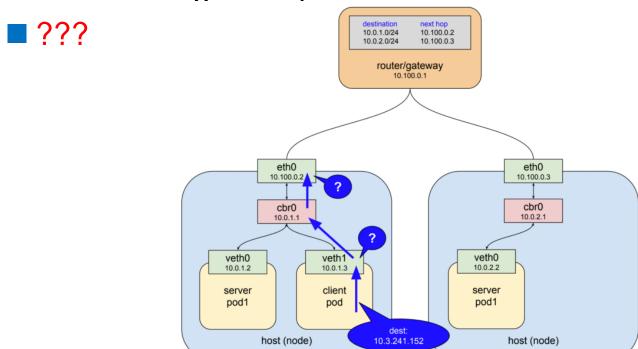


- A service
 - A type of kubernetes resource
 - Having a configured proxy that can forward requests to a set of pods
- Service's IP
 - Belong to "Service Network"
 - Different with pod's network and node's network
 - Services needed their own, stable, non-conflicting network address space
 - A system of virtual IPs is not stable
 - The service network does not exist

??? How could the request reach the running pod via service's IP ???

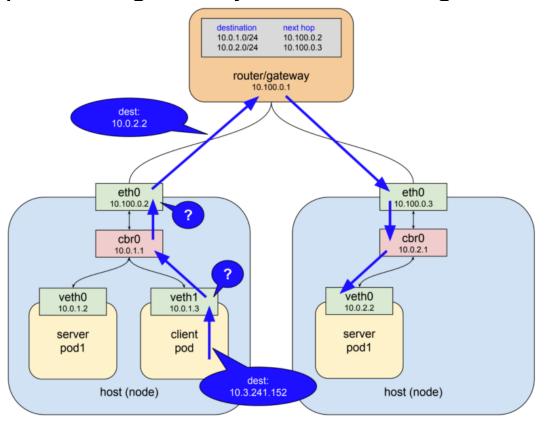


- The client make an HTTP request to the service via DNS name
- The DNS system resolves service's name to cluster IP
- The client creates HTTP request to cluster IP
- The router/gateway does not know about cluster IP





Strategy: upstream gateway for forwarding unknown address



- How do those packets popped out at the right place in cluster
- => The magician here called kube-proxy



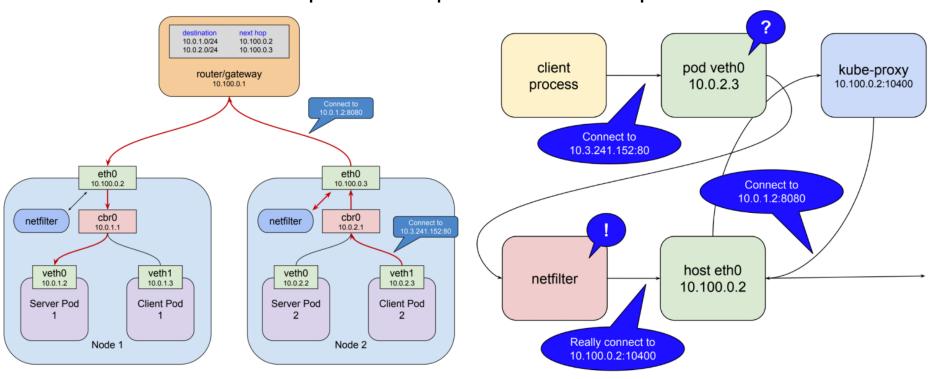
- What is kube-proxy
 - A kubernetes resource
 - A record in a central database
 - Quite different from a typical reverse-proxy
- Issue: How can we listen on a port or open a connection through an interface that doesn't exist???
 - The host's ethernet interface
 - The virtual ethernet interfaces on the pod network

Solution

- Using netfilter as a kernel space proxy
 - Running in kernel space and gets a look at every packet at various points
 - Redirecting the packet to another destination

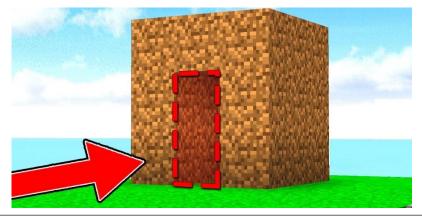


- Kube-proxy
 - Open a port (10400) on the node interface to listen for requests to the registered service
 - Inserts netfilter rules to reroute packets destined for the service IP to its own port
 - Forwards those requests to a pod on a service port





- Is the services proxy system durable?
 - Running as a system unit => it will be restarted if it fails
- Is the service proxy aware of healthy server pods?
 - kube-proxy listens to the master apiserver for changes in the cluster (service, endpoint)
 - Service creation => kube-proxy => create the necessary rules
 - Serivce deletion => kube-proxy => remove the related rules
 - Health-check are performed by kubelet
 - Unhealthy endpoints => apiserver => kube-proxy => netfilter rules edition
- Only working for request that originate from inside the cluster





How clients *outside* the cluster can connect to pods using the same service network

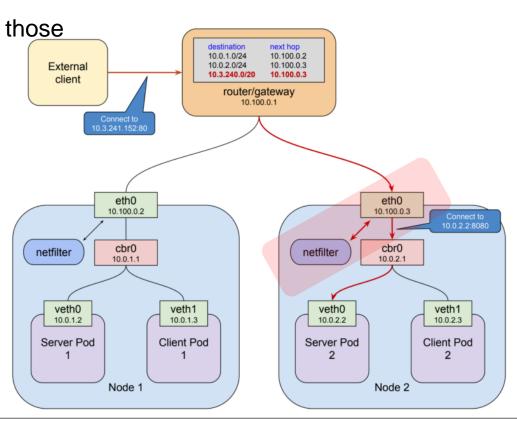
INGRESS



- Try to make use of the routing infrastructure
 - Calling ClusterIP and Port
 - Problem: The ClusterIP only reachable from node interface
 - ⇒How can we forward traffic from a public IP to an IP that is only reachable from node?

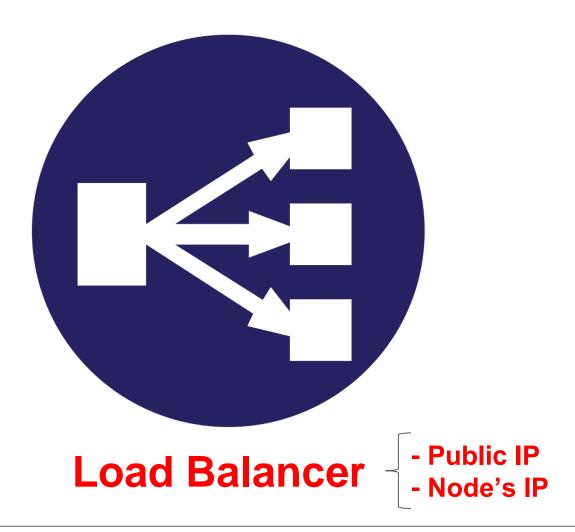
⇒Should we add a route to get those packets to one of the nodes?

- It works
- Not stable
- Not optimal



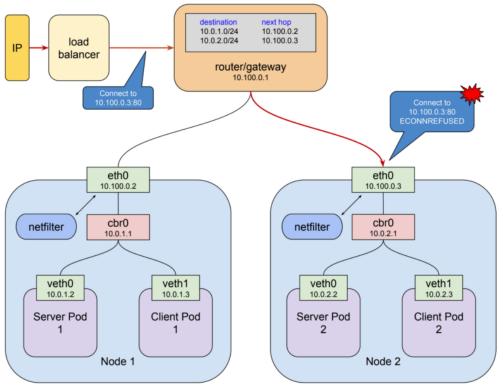


Solution





No process listening on 10.100.0.3:80



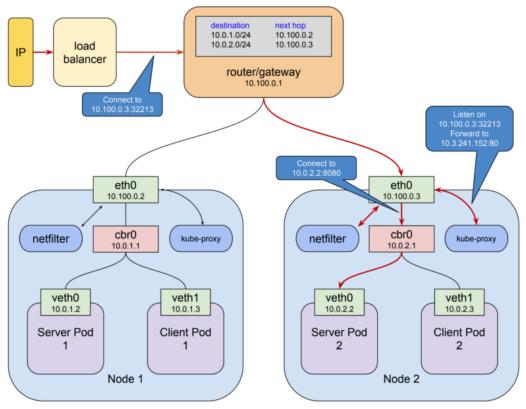
- netfilter rules don't match destination address
 - Service network: match rule, but not routable
 - Node network: routable, but not match rule



Bridge



- NodePort (ClusterIP extension)
 - Reachable at IP address of Node
 - Reachable in Service Network
 - How it works?
 - Kube-proxy allocates a port in 30000-32767
 - Open this port on the eth0 of every node
 - Connection to this port are forwarded to the service's cluster IP



Issues:

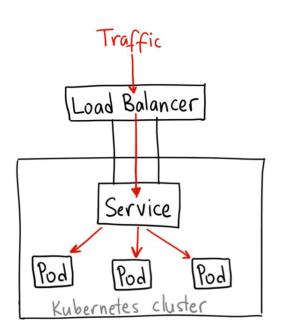
- Exposing Port is a non-standard port
- Limited resource (2768 ports)

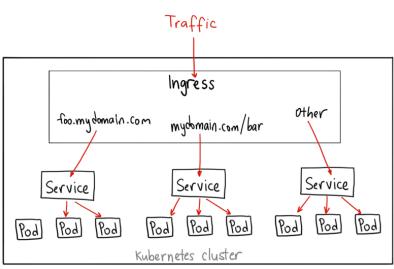


- LoadBalancer Service (NodePort Extension)
 - Build out a complete ingress path with external IP

```
$ kubectl get svc service-test
NAME     CLUSTER-IP     EXTERNAL-IP     PORT(S)     AGE
openvpn     10.3.241.52     35.184.97.156     80:32213/TCP 5m
```

- If service isn't destroyed and recreated, the IP won't change.
- Limitation
 - Support single service (cannot use single load balancer to multiple service)
- Ingress
 - Sit in front of multiple services
 - Act as an entrypoint into the cluster.





v2.0 is coming







"Kubernetes networking with Istio, Envoy, Cilium"

Q & A